

并行算法试题答题卷

林宇健 2018202296

2019 年 2 月 22 日

1 并行矩阵向量乘法

1.1 问题描述和分析

编程计算 Ax ，其中 A 是 $m \times n$ 的稠密矩阵， x 是 n 维列向量，分别采用1, 4, 8, 16台处理机计算。给出并行算法，及并行效率分析。

记 $y = Ax$ 。为简便起见，取 A 是 $n \times n$ 的方阵， $n = 2048$ 可以整除1, 4, 8, 16，从而保证每个进程储存的向量块维度相同。对于不能整除4, 8, 16的 n ，可通过循环存储等方式为各个进程分配矩阵和向量的数据。在计算时编程随机生成了 $matrix_{2048 \times 2048}$ 和 $vector_{2048 \times 1}$ 作为待计算的矩阵和向量（计算程序略去）。

我们采用一维行划分的方式并行计算矩阵向量乘法。假设矩阵 A 按逐行一维块划分为 p 个块（ p 表示进程数），即 $A = [A_1, A_2, \dots, A_p]^T$, $A_k = [A_{k,0}, A_{k,1}, \dots, A_{k,p}]$ 。其中

$$A_{k,j} = \begin{bmatrix} a_{k \times n/p+1, j \times n/p+1} & a_{k \times n/p+1, j \times n/p+2} & \cdots & a_{k \times n/p+1, j \times n/p+n/p} \\ a_{k \times n/p+2, j \times n/p+1} & a_{k \times n/p+2, j \times n/p+2} & \cdots & a_{k \times n/p+2, j \times n/p+n/p} \\ \vdots & \vdots & \ddots & \vdots \\ a_{k \times n/p+n/p, j \times n/p+1} & a_{k \times n/p+n/p, j \times n/p+2} & \cdots & a_{k \times n/p+n/p, j \times n/p+n/p} \end{bmatrix}$$

与之相对应，向量 x 和向量 y 也分为 p 个块，其第 k ($0 \leq k \leq p-1$)个块分别为

A				X	
$A_{1,1}$	$A_{1,2}$	\cdots	$A_{1,p}$	x_1	P_1
$A_{2,1}$	$A_{2,2}$	\cdots	$A_{2,p}$	x_2	P_2
\vdots	\vdots	\ddots	\vdots	\vdots	\vdots
$A_{p,1}$	$A_{p,2}$	\cdots	$A_{p,p}$	x_p	P_p

```
1 !*****
2 !
3 ! matrix_mul_vector_parallel.f90
4 ! 并行矩阵向量乘法，基于行划分方法进行并行化
5 !
6 !*****
7
```

```

8  program parallel_Mat_mul_Vec
9
10  use mpi
11  implicit none
12
13  integer, parameter :: N = 2048
14  integer :: my_left, my_right
15  integer :: IERR, NPROC, NSTATUS(MPI_STATUS_SIZE)
16  integer :: myrank, myleft, myright, myfile, buf_size, cnt
17  real(4) :: startwtime, endwtime, wtime
18  real(4), allocatable :: matrix(:, :), vector(:, :), answer(:, :)
19  real(4), allocatable :: matrix_buf(:, :), vector_buf(:, :), answer_buf(:, :)
20  character(len = 2) :: sTemp
21
22  call cpu_time(startwtime)
23
24  call mpi_init(IERR)
25  call mpi_comm_rank(MPI_COMM_WORLD, myrank, IERR)
26  call mpi_comm_size(MPI_COMM_WORLD, NPROC, IERR)
27
28  buf_size = N / NPROC
29  myleft = my_left(myrank, NPROC)
30  myright = my_right(myrank, NPROC)
31
32  allocate(matrix_buf(N, buf_size)) ! Fortran的矩阵储存方式为列储存, 需要
33                                     ! 进行一次转置, 因此设置读取缓存空间
34  allocate(vector_buf(buf_size, 1)) ! 接收其他进程储存的向量所需要的缓存空间
35  allocate(matrix(buf_size, N))
36  allocate(vector(buf_size, 1))
37  allocate(answer(N, 1))
38  allocate(answer_buf(N, 1))
39
40  ! 读取矩阵
41  call mpi_file_open(MPI_COMM_WORLD, "matrix", MPI_MODE_RDONLY, MPI_INFO_NULL, &
42  & myfile, IERR)
43  call mpi_file_seek(myfile, myrank*N*buf_size*sizeof(MPI_REAL), MPI_SEEK_SET, &
44  & IERR)
45  call mpi_file_read(myfile, matrix_buf, N*buf_size, MPI_REAL, NSTATUS, IERR)
46  call mpi_file_close(myfile, IERR)
47  matrix = transpose(matrix_buf)
48
49  ! 读取向量
50  call mpi_file_open(MPI_COMM_WORLD, "vector", MPI_MODE_RDONLY, MPI_INFO_NULL, &
51  & myfile, IERR)
52  call mpi_file_seek(myfile, myrank*buf_size*sizeof(MPI_REAL), MPI_SEEK_SET, IERR)
53  call mpi_file_read(myfile, vector, buf_size, MPI_REAL, NSTATUS, IERR)
54  call mpi_file_close(myfile, IERR)
55
56  answer = 0
57  deallocate(matrix_buf) ! 释放矩阵缓存空间用于储存每一次计算时的矩阵块
58  allocate(matrix_buf(buf_size, buf_size))
59  ! 循环进程中储存的所有矩阵块
60  do cnt = 0, NPROC
61      ! 计算对应矩阵块与向量的乘积
62      matrix_buf = matrix(:, mod(myrank+cnt, NPROC)*buf_size+1:(mod(myrank+cnt, NPROC) &
63  & +1)*buf_size)
64      answer(myrank*buf_size+1:(myrank+1)*buf_size, :) = matmul(matrix_buf, vector) &
65  & + answer(myrank*buf_size+1:(myrank+1)*buf_size, :)
66      ! 进行一次向量块的传递(向上)
67      call mpi_send(vector, buf_size, MPI_REAL, myleft, myrank, MPI_COMM_WORLD, IERR)
68      call mpi_recv(vector_buf, buf_size, MPI_REAL, myright, myright, &
69  & MPI_COMM_WORLD, NSTATUS, IERR)
70      vector = vector_buf
71  end do
72
73  ! 全规约结果向量, 并行输出到文件
74  call mpi_allreduce(answer, answer_buf, N, MPI_REAL, MPI_SUM, MPI_COMM_WORLD, IERR)
75  call mpi_file_open(MPI_COMM_WORLD, "answer", MPI_MODE_CREATE+MPI_MODE_WRONLY, &
76  & MPI_INFO_NULL, myfile, IERR)
77  call mpi_file_seek(myfile, myrank*buf_size*sizeof(MPI_REAL), MPI_SEEK_SET, IERR)
78  call mpi_file_write(myfile, answer_buf(myrank*buf_size+1, 1), buf_size, MPI_REAL, &
79  & MPI_STATUS_IGNORE, IERR)
80  call mpi_file_close(myfile, IERR)
81
82  ! 将各进程的运行时间记录到文件中
83  call cpu_time(endwtime)
84  wtime = (endwtime - startwtime) * 1000
85  write(sTemp, '(i2)') NPROC
86  call mpi_file_open(MPI_COMM_WORLD, "walltime"//trim(adjustl(sTemp)), MPI_MODE_CREATE &
87  & +MPI_MODE_WRONLY, MPI_INFO_NULL, myfile, IERR)
88  call mpi_file_seek(myfile, myrank*sizeof(MPI_REAL), MPI_SEEK_SET, IERR)

```

```

89      call mpi_file_write(myfile, wtime, 1, MPI_REAL, MPI_STATUS_IGNORE, IERR)
90      call mpi_file_close(myfile, IERR)
91
92      deallocate(matrix)
93      deallocate(vector)
94      deallocate(answer)
95      deallocate(matrix_buf)
96      deallocate(vector_buf)
97      deallocate(answer_buf)
98      call mpi_finalize(IERR)
99
100 end program parallel_Mat_mul_Vec
101
102
103 !-----子程序与函数部分-----
104 integer function my_left(myrank, nproc) result(ans)
105
106     implicit none
107     integer, intent(in) :: myrank, nproc
108
109     ans = myrank - 1
110     if (0 == myrank) ans = nproc - 1
111
112 end function my_left
113
114
115 integer function my_right(myrank, nproc) result(ans)
116
117     implicit none
118     integer, intent(in) :: myrank, nproc
119
120     ans = myrank + 1
121     if (nproc-1 == myrank) ans = 0
122
123 end function my_right

```

```

1  !*****
2  !
3  ! matrix_mul_vector_serial.f90
4  ! 串行矩阵向量乘法
5  !
6  !*****
7
8 program serial_Mat_mul_Vec
9
10     implicit none
11
12     integer, parameter :: N = 2048
13     integer :: i, j
14     real(4) :: startwtime, endwtime
15     real(4) :: matrix(N, N), vector(N, 1), answer(N, 1) = 0
16
17     call cpu_time(startwtime)
18
19     ! 读取矩阵
20     open(10, file = 'matrix', access = 'direct', form = 'unformatted', recl = 4*N*N)
21     read(10, rec = 1) ((matrix(i,j), j = 1, N), i = 1, N)
22     close(10)
23     matrix = transpose(matrix) ! Fortran的矩阵储存方式为列储存, 需要
24                               ! 进行一次转置
25
26     ! 读取向量
27     open(20, file = 'vector', access = 'direct', form = 'unformatted', recl = 4*N)
28     read(20, rec = 1) (vector(i, 1), i = 1, N)
29     close(20)
30
31     answer = matmul(matrix, vector)
32
33     ! 输出结果到向量文件
34     open(30, file = 'answer', access = 'direct', form = 'unformatted', recl = 4)
35     do i = 1, N
36         write(30, rec = i) answer(i, 1)
37     end do
38
39     call cpu_time(endwtime)
40     open(40, file = 'walltime', access = 'direct', form = 'unformatted', recl = 4)
41     write(40, rec = 1) (endwtime - startwtime) * 1000
42     close(40)
43
44 end program serial_Mat_mul_Vec

```

```

1  #!/bin/bash
2  # nohup sh autoexec.sh > /dev/null 2>&1 &
3
4  gfortran random_matrix.f90 -o matrix.out
5  gfortran random_vector.f90 -o vector.out
6  ./matrix.out
7  ./vector.out
8  mpif90 matrix_mul_vector_parallel.f90
9  mpirun -np 1 a.out
10 mpirun -np 4 a.out
11 mpirun -np 8 a.out
12 mpirun -np 16 a.out
13 gfortran matrix_mul_vector_serial.f90 -o b.out
14 ./b.out
15 mpif90 matrix_mul_vector_serial_with_mpi.f90 -o c.out
16 mpirun -np 1 c.out
17 gfortran walltime.for -o d.out
18 ./d.out

```

```

1  C*****
2  C
3  C    walltime.for
4  C    计算串行程序和并行程序的运行时间并显
5  C
6  C*****
7
8  PROGRAM WALLTIME_FOR
9
10 REAL*4 T, TM, T1, T4(4), T8(8), T16(16)
11
12 OPEN(8, FILE = 'walltime', ACCESS = 'DIRECT', FORM = 'UNFORMATTED
13 & ', RECL = 4)
14 READ(8, REC = 1) T
15 CLOSE(8)
16 PRINT *, "serial program walltime: ", T, "ms"
17
18 OPEN(8, FILE = 'walltime_mpiio', ACCESS = 'DIRECT',
19 & FORM = 'UNFORMATTED', RECL = 4)
20 READ(8, REC = 1) TM
21 CLOSE(8)
22 PRINT *, "serial program(MPI I/O) walltime: ", TM, "ms"
23
24 OPEN(8, FILE = 'walltime1', ACCESS = 'DIRECT', FORM = 'UNFORMATTED
25 & ', RECL = 4)
26 READ(8, REC = 1) T1
27 CLOSE(8)
28 PRINT *, "parallel program walltime (1 process): ", T1, "ms"
29
30 OPEN(8, FILE = 'walltime4', ACCESS = 'DIRECT', FORM = 'UNFORMATTED
31 & ', RECL = 4*4)
32 READ(8, REC = 1) T4
33 CLOSE(8)
34 SUM4 = .0
35 DO 40 I = 1, 4
36     SUM4 = SUM4 + T4(I)
37 40 CONTINUE
38 PRINT *, "parallel program walltime (4 process): ", SUM4/4, "ms"
39
40 OPEN(8, FILE = 'walltime8', ACCESS = 'DIRECT', FORM = 'UNFORMATTED
41 & ', RECL = 4*8)
42 READ(8, REC = 1) T8
43 CLOSE(8)
44 SUM8 = .0
45 DO 80 I = 1, 8
46     SUM8 = SUM8 + T8(I)
47 80 CONTINUE
48 PRINT *, "parallel program walltime (8 process): ", SUM8/8, "ms"
49
50 OPEN(8, FILE = 'walltime16', ACCESS = 'DIRECT', FORM =
51 & 'UNFORMATTED', RECL = 4*16)
52 READ(8, REC = 1) T16
53 CLOSE(8)
54 SUM16 = .0
55 DO 160 I = 1, 16
56     SUM16 = SUM16 + T16(I)
57 160 CONTINUE
58 PRINT *, "parallel program walltime (16 process): ", SUM16/16,
59 & "ms"
60

```

61 | `END PROGRAM WALLTIME_FOR`
